


Do the action channels separate failures any better? (logreg)

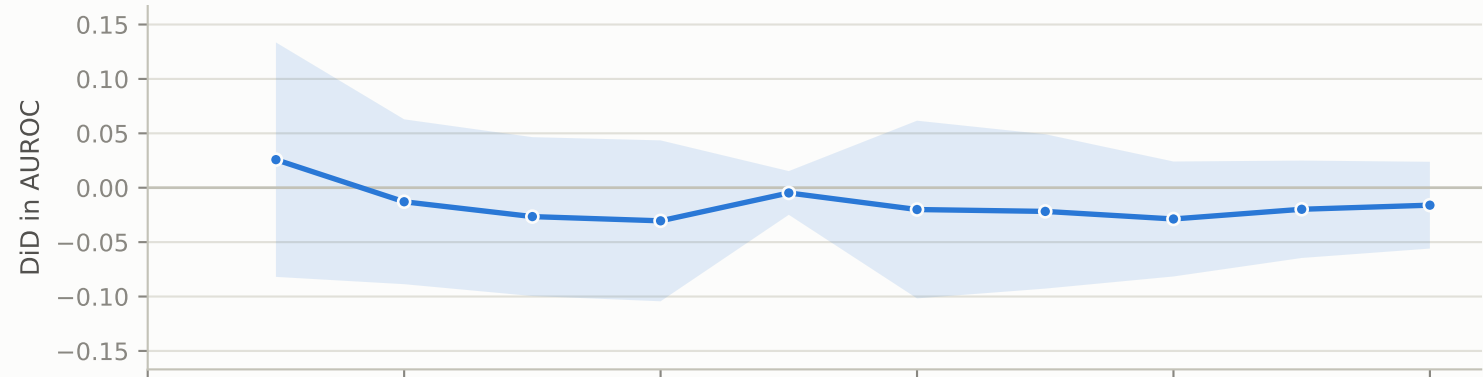
Q3, difference-in-differences: $(\text{failure} - \text{success} \mid \text{obs} + \text{act}) - (\text{failure} - \text{success} \mid \text{obs})$, formed within each CV fold and averaged; the band is a 95% t interval over the folds. $y = 0$ is the observation-only baseline of plot 1, and the success curve controls for calibration-vs-test shift unrelated to failure.

push_t


above 0: the action channels widen the failure-minus-success gap




pretzel




push_chair



sorting



stacking



0.0 0.2 0.4 0.6 0.8 1.0
normalized episode time t (the pool is every step up to t)