




Do the action channels help? Difference-in-differences vs the observation-only baseline (Q3)
DiD = (fail – succ | obs+act) – (fail – succ | obs), in AUROC points, formed per fold
band = 95% CI (Student-t, 4 df); right axis rescales it to % of that task's obs-only gap
Push T — nominal 50 ep, test 102 success / 198 failure, d_obs = 64



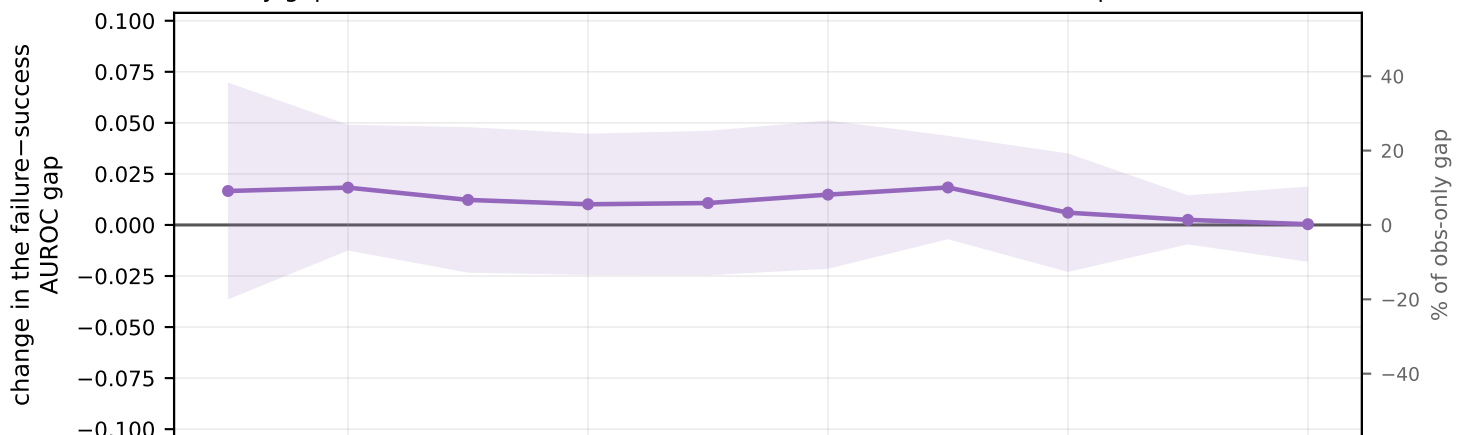
Pretzel — nominal 10 ep, test 10 success / 10 failure, d_obs = 512 + d_act = 1440
obs-only gap +0.218; mean DiD +0.034 = +16% of it; CI excludes 0 at 4/10 points



Push Chair — nominal 10 ep, test 10 success / 10 failure, d_obs = 96



Sorting — nominal 50 ep, test 255 success / 145 failure, d_obs = 128 + d_act = 768
obs-only gap +0.182; mean DiD +0.011 = +6% of it; CI excludes 0 at 0/10 points



Stacking — nominal 50 ep, test 315 success / 485 failure, d_obs = 640 + d_act = 768
obs-only gap +0.343; mean DiD -0.016 = -5% of it; CI excludes 0 at 2/10 points

